

Your notebook (host)

Docker (host network)

Python code

py4j

JVM: the DRIVER

- plans each query
- splits it into tasks
- serves jar files

① need resources

spark-master

② starts executors

④ tasks

③ connect back

⑤ results

spark-worker-1, -2
executors (JVM)
+ Python 3.13 for UDFs

